

Games That Talk: New Foundations for Composable Security

Lúcás Críostóir Meier
lucas@cronokirby.com

December 8, 2022

Abstract

We do things with UC security.

1 Introduction

[Mei22]

Definition 1.1 (Adversaries). An adversary is a cool thing.

Theorem 1.1 (Cool Beans). Woah mama

And that's what matters.

■

Lemma 1.2. Woah mama again!

Corollary 1.3. Woah mama again!

Γ^0
$x \leftarrow 3$
if $x + 2$:
$y \xleftarrow{\$} \mathbb{F}_q$
$m \Rightarrow \langle \mathcal{P}_i, \mathcal{P}_j \rangle$ $y \leftarrow 4$
$m \Leftarrow \langle \text{OT}, \mathcal{P}_i \rangle$ $x \leftarrow 3$
$\frac{\text{Foo}(x, y):}{\text{Bar}(x, y)}$

Game 1.1: Some Game

1.1 Relevance of Time Travel

stuff

IND-CCA

$x \leftarrow 4$

Protocol 1.2: Some Protocol

IND-CCA

$x \leftarrow 4$

Protocol 1.3: Some Protocol

IND-CCA

$x \leftarrow 3$

Functionality 1.4: Encryption

2 State-Separable Proofs

3 Games That Talk

3.1 Async Functions

While the intuition of yield statements is simple, defining them precisely is a bit more tricky.

Definition 3.1 (Yield Statements). We define the semantics of **yield** by compiling functions with such statements to functions without them.

Note that we don't define the semantics for functions which still contain references to oracles. Like before, we can delay the definition of semantics until all of the pseudo-code has been inlined.

A first small change is to make it so that the function accepts one argument, a binary string, and all yield points also accept binary strings as continuation. Like with plain packages, we can implement richer types on top by adding additional checks to the well-formedness of binary strings, aborting otherwise.

The next step is to make it so that all the local variables of the function F are present in the global state. So, if a local variable v is present, then every use of v is replaced with a use of the global variable $F.v$ in the package. This allows the state of the function to be saved across yields.

The next step is transforming all the control flow of a function to use **ifgoto**, rather than structured programming constructs like **while** or **if**. The function is

broken into lines, each of which contains a single statement. Each line is given a number, starting at 0. The execution of a function F involves a special variable pc , representing the current line being executing. Excluding **yield** and **return** a single line statement has one of the forms:

$$\begin{aligned}\langle \text{var} \rangle &\leftarrow \langle \text{expr} \rangle \\ \langle \text{var} \rangle &\overset{\$}{\leftarrow} \langle \text{dist} \rangle\end{aligned}$$

which have well defined semantics already. Additionally, after these statements, we set $pc \leftarrow pc + 1$.

The semantics of **ifgoto** $\langle \text{expr} \rangle i$ is:

$$pc \leftarrow \text{if } \langle \text{expr} \rangle \text{ then } i \text{ else } pc + 1$$

This gives us a conditional jump, and by using **true** as the condition, we get a standard unconditional jump.

This allows us to define **if** and **while** statements in the natural way.

Finally, we need to augment functions to handle **yield** and **return** statements. To handle this, each function F also has an associated variable $F.pc$, which stores the program counter for the function. This is different than the local pc which is while the function is execution. $F.pc$ is simply used to remember the program counter after a yield statement.

The function now starts with:

$$\text{ifgoto true } F.pc$$

This has the effect of resuming execution at the saved program counter.

Furthermore, the input variable x to F is replaced with a special variable **input**, which holds the input supplied to the function. At the start of the function body, we add:

$$0 : F.x \leftarrow \text{input}$$

to capture the fact that the original input variable needs to get assigned to the **input** to the function.

The semantics of $F.m \leftarrow \text{yield } v$ are:

$$\begin{aligned}(i - 1) : F.pc &\leftarrow i + 1 \\ i : &\text{return } (\text{yield}, v) \\ (i + 1) : F.m &\leftarrow \text{input}\end{aligned}$$

The semantics of **return** v become:

$$\begin{aligned}F.pc &\leftarrow 0 \\ \text{return } &(\text{return}, v)\end{aligned}$$

The main difference is that we annotate the return value to be different than yield statements, but otherwise the semantics are the same.

□

Note that while calling a function which can yield will notify the caller as to whether or not the return value was *yielded* or *returned*, syntactically the caller often ignores this, simply doing $x \leftarrow F(\dots)$, meaning that they simply use return value x , discarding the tag.

Syntax 3.2. In many cases, no value is yielded, or returned back, which we can write as:

yield

which is shorthand for:

• \leftarrow **yield** •

i.e. just yielding a dummy value and ignoring the result.

□

In such situations, often we don't particularly care about the intermediate yields of a function, and want to wait for the final result, potentially yielding to our own caller. We define these semantics via the **await** statement.

Syntax 3.3 (Await Statements). We define the semantics of $v \leftarrow \mathbf{await} F(\dots)$ in a straightforward way:

$$\begin{aligned} &(\text{tag}, v) \leftarrow (\text{yield}, \perp) \\ &\mathbf{while} \text{ tag} = \text{yield} : \\ &\quad \mathbf{if} \ v \neq \perp : \\ &\quad \quad \mathbf{yield} \\ &\quad (\text{tag}, v) \leftarrow F(\dots) \end{aligned}$$

In other words, we keep calling the function until it actually returns its final value, but we do yield to our caller whenever our function yield, but we do yield to our caller whenever our function yields.

□

Sometimes we want to await several values at once, returning the first one which completes. To that end, we define the **select** statement.

Syntax 3.4 (Select Statements). Select statements generalize await statements in that they allow waiting for multiple events concurrently.

More formally, we define:

```

select :
   $v_1 \leftarrow \mathbf{await} F_1(\dots) :$ 
     $\langle \text{body}_1 \rangle$ 
   $\vdots$ 
   $v_n \leftarrow \mathbf{await} F_n(\dots) :$ 
     $\langle \text{body}_n \rangle$ 

```

As follows:

```

 $(\text{tag}_i, v_i) \leftarrow (\text{yield}, \perp)$ 
 $i \leftarrow 0$ 
while  $\nexists i. \text{tag}_i \neq \text{yield} :$ 
  if  $i \geq n :$ 
     $i \leftarrow 0$ 
  yield
   $i \leftarrow i + 1$ 
   $(\text{tag}_i, v_i) \leftarrow F_i(\dots)$ 
   $\langle \text{body}_i \rangle$ 

```

Note that the order in which we call the functions is completely deterministic, and fair. It's also important that we yield, like with await statements, but we only do so after having pinged each of our underlying functions at least once. This is so that if one of the function is immediately ready, we never yield.

□

3.2 Channels and System Composition

Definition 3.5 (Systems). A *system* is a package which uses channels.

We denote by $\text{InChan}(S)$ the set of channels the system receives on, and $\text{OutChan}(S)$ the set of channels the system sends on, and define

$$\text{Chan}(S) := \text{OutChan}(S) \cup \text{InChan}(S)$$

□

Definition 3.6. We can compile systems to not use channels. We denote by $\text{NoChan}(S)$ the package corresponding to a system S , with the use of channels replaced with function calls.

Channels define two new syntactic constructions, for sending and receiving along a channel. We replace these with function calls as follows:

Sending, with $m \Rightarrow P$ becomes:

`Channels.Send(m, P)`

Receiving, with $m \Leftarrow P$ becomes:

$m \leftarrow \mathbf{await}$ `Channels.Recv(P)`

Receiving is an asynchronous function, because the channel might not have any available messages for us.

□

```
Channels( $\{A_1, \dots, A_n\}$ )
 $q[A_i] \leftarrow \text{FifoQueue.New}()$ 
Send( $m, A_i$ ):
   $q[A_i].\text{Push}(m)$ 
Recv( $A_i$ ):
  while  $q[A_i].\text{IsEmpty}()$ 
  yield
   $q[A_i].\text{Next}()$ 
```

Game 3.1: Channels

Armed with the syntax sugar for channels, and the Channels game, we can convert a system S into a package via:

$$\text{SysPack}(S) := \text{NoChan}(S) \circ (\text{Channels}(\text{Chan}(S)) \otimes 1(\text{In}(S)))$$

This package will have the same input and output functions as the system S , but with the usage of channels replaced with actual semantics.

This allows us to lift our standard equality relations on packages onto *systems*.

Definition 3.7. Given some equality relation \sim on packages, we can lift that relation to systems by defining:

$$A \sim B \iff \text{SysPack}(A) \sim \text{SysPack}(B)$$

□

Definition 3.8 (System Tensoring). Given two systems, A and B , we can define their tensor product $A * B$, which is any system satisfying:

$$\text{SysPack}(A * B) = (\text{NoChan}(A) \otimes \text{NoChan}(B)) \circ \left(\begin{array}{c} \text{Channels}(\text{Chan}(A) \cup \text{Chan}(B)) \\ \otimes \\ 1(\text{In}(A) \cup \text{In}(B)) \end{array} \right)$$

□

Lemma 3.1. System tensoring is associative, i.e. $A * (B * C) = (A * B) * C$.

Proof: This follows from the associativity of \otimes and \cup . ■

Definition 3.9 (Overlapping Systems). Two systems A and B overlap if $\text{Chan}(A) \cap \text{Chan}(B) \neq \emptyset$.

In the case of non-overlapping systems, we write $A \otimes B$ instead of $A * B$, insisting on the fact that they don't communicate.

Definition 3.10 (System Composition). Given two systems, A and B , we can define their (horizontal) composition $A \circ B$ as any system satisfying: **add constraints like for packages**

$$\text{SysPack}(A \circ B) = \text{SysPack}(A) \circ \text{SysPack}(B)$$

□

Lemma 3.2. System composition is associative, i.e. $A \circ (B \circ C) = (A \circ B) \circ C$.

Proof: This follows from the associativity of \circ for *packages*. ■

Definition 3.11 (System Games). Analogously to games, we define a *system game* as a pair of systems S_b with $\text{In}(S_b) = \emptyset$.

Definition 3.12 (System Game Reductions). We can also define notions of reductions for system games.

First, we define:

$$\epsilon(\mathcal{A} \circ S_b) := \epsilon(\mathcal{A} \circ \text{SysPack}(S_b))$$

We then also use the syntax sugar of:

$$S_b \leq f(G_b^1, G_b^2, \dots)$$

as shorthand for, $\forall \mathcal{A}. \exists \mathcal{B}_1, \dots$:

$$\epsilon(\mathcal{A} \circ S_b) \leq f(\epsilon(\mathcal{B}_1 \circ G_b^1), \epsilon(\mathcal{B}_2 \circ G_b^2), \dots)$$

We also sometimes omit explicitly writing S_b , instead writing just S , if it's clear that we're talking about a pair of systems.

□

Similar properties hold for reductions:

Lemma 3.3. $A \circ G_b \leq G_b$.

Proof: $\text{SysPack}(A \circ G_b) = \text{SysPack}(A) \circ \text{SysPack}(G_b) \leq \text{SysPack}(G_b)$. ■

4 Protocols and Composition

5 Differences with UC Security

6 Examples

7 Further Work

8 Conclusion

References

- [Mei22] Lúcas Crístóir Meier. MPC for group reconstruction circuits. Cryptology ePrint Archive, Report 2022/821, 2022. <https://eprint.iacr.org/2022/821>.